

Theoretical Understanding

Q1: Explain the primary differences between TensorFlow and PyTorch. When would you choose one over the other?

Answer:

The primary difference is in how computation graphs are handled: TensorFlow (especially v1) uses static graphs, while PyTorch uses dynamic (eager) execution, making it more intuitive and Pythonic.

- Choose PyTorch for research, experimentation, or education, due to its flexibility and ease of debugging.
 - Choose TensorFlow for production deployment, especially on mobile or in distributed systems, due to tools like TensorFlow Serving and TensorFlow Lite.
-

Q2: Describe two use cases for Jupyter Notebooks in AI development.

Answer:

1. Prototyping and Experimentation: Jupyter allows for quick testing and visualization of models, data preprocessing steps, and evaluation metrics interactively.
 2. Data Exploration and Visualization: It's ideal for analyzing datasets with graphs, tables, and plots using libraries like `matplotlib`, `seaborn`, and `pandas`.
-

Q3: How does spaCy enhance NLP tasks compared to basic Python string operations?

Answer:

spaCy provides linguistically informed processing like tokenization, POS tagging, named entity recognition, and dependency parsing, which go far beyond basic string methods (e.g., `.split()` or `.find()`). This allows for accurate and contextual understanding of text, essential for building robust NLP applications.

Comparing scikit-learn, tensorflow and pytorch

1. Architecture

Feature	TensorFlow	PyTorch	Scikit-learn
Core Design	Static computation graph (eager by default in TF 2.x)	Dynamic computation graph (define-by-run)	No computation graph; traditional ML pipeline
Language Support	Python, C++, JavaScript (via TensorFlow.js)	Primarily Python (C++ backend)	Pure Python
Deployment	TensorFlow Serving, TFLite, TensorFlow.js	TorchScript, ONNX	Not focused on deployment
Hardware Support	CPU, GPU, TPU	CPU, GPU	CPU (limited GPU support via joblib/sklearnex)

2. Syntax & Programming Style

Feature	TensorFlow	PyTorch	Scikit-learn
Ease of Use	Moderate to steep learning curve	Pythonic, intuitive	Very beginner-friendly
Model Definition	Using <code>tf.keras.Model</code> , <code>Sequential</code> , functional API	Use <code>torch.nn.Module</code> and forward method	Simple API (<code>fit</code> , <code>predict</code> , <code>score</code>)
Debugging	Harder in TF 1.x, better in TF 2.x (eager mode)	Easier due to dynamic graph	Straightforward
Training Loop	Can use <code>model.fit()</code> or custom loops	Typically custom loops	Handled internally

3. Use Cases

Use Case	TensorFlow	PyTorch	Scikit-learn
Deep Learning (CNNs, RNNs)	✅ Industry standard	✅ Preferred in research	❌ Not suitable
Production/Deployment	✅ Optimized for production	⚠️ Possible but less mature	❌ Not designed for production
Research & Experimentation	⚠️ Good with TF 2.x	✅ Excellent	⚠️ Limited to classical ML
Classical Machine Learning	⚠️ Some support via <code>tf.estimator</code>	⚠️ Basic support	✅ Full support (SVMs, Trees, Clustering, etc.)
Mobile/Embedded AI	✅ via TFLite	⚠️ Possible with optimization	❌

4. When to Use Which

Scenario	Use This Tool	Why
You want to build a deep learning model for production use	TensorFlow	Strong deployment tools (TFLite, TensorFlow Serving)
You're doing research or prototyping with deep learning	PyTorch	Flexible, intuitive, dynamic graph model
You're building classical ML models (e.g., SVMs, Random Forests, KNN)	Scikit-learn	Simple, mature, and optimized for classical ML workflows
You want to deploy AI to mobile devices	TensorFlow (TFLite)	Excellent support for on-device inference
You're teaching or learning ML basics	Scikit-learn	Clean, minimal code with well-documented algorithms
You need explainable ML with small datasets	Scikit-learn	Easier to explain and visualize decision boundaries
You want to convert models to ONNX for interoperability	PyTorch or TensorFlow	Both support ONNX export with some effort

Summary Table

Feature	TensorFlow	PyTorch	Scikit-learn
Best For	Production DL	Research DL	Classical ML
Learning Curve	Moderate/Steep	Moderate	Easy
Model Types	Deep Learning	Deep Learning	Traditional ML
Production Support	Excellent	Good (improving)	Minimal
GPU Support	Yes	Yes	Limited

Results/Outputs

Ethical Considerations

Potential Biases

In MNIST:

- **Class Imbalance:** Some digits may be underrepresented, leading to lower accuracy for those classes.
- **Overfitting on writing style:** The model might generalize poorly to handwriting styles not present in the training data.

In Amazon Reviews:

- **Sentiment bias:** The model may learn bias from skewed review distributions (e.g., more positive reviews).
- **Demographic bias:** Reviews might reflect cultural, gender, or product-related biases that are not explicit in the data.